



Summary

Modules that compete for action choices
have the promise of good scaling behavior

Agent “Walter” walks down a sidewalk, avoiding blue obstacles and picking up purple litter



Sidewalk




Obstacle




Litter

State Spaces definitions




Obstacle & Litter: heading and distance to target


Sidewalk: shortest distance to sidewalk center




State defined by nearest target



Policy




Module for Litter Cleanup



Heading from agent's perspective

V is value of Policy

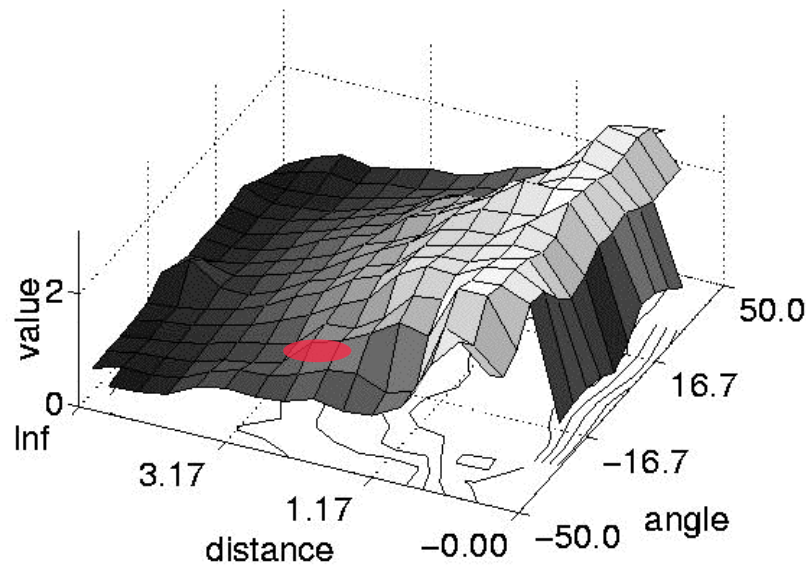


$$V(s) = \max_a Q(s,a)$$


Learned Module Behaviors

Red dot indicates each module's state





Litter



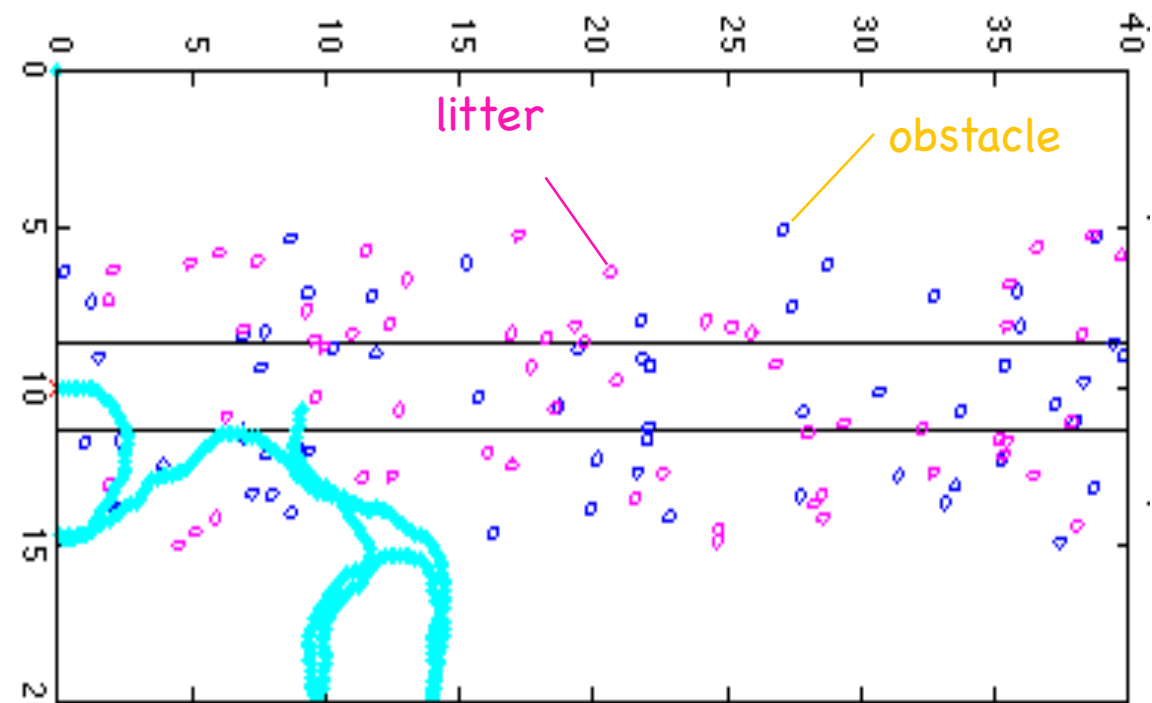
Sidewalk




Obstacles




Overhead view of trajectories showing stages in learning



Initial
performance



After
100 episodes



After
150 episodes